

# Распознавание рукописных цифр из датасета MNIST с помощью сверточной нейросети.

immediate

December 14, 2021

**Выполнила студент 431 группы Алтынова Анна.**

**Цель:** построить классификатор для рукописных цифр на основе нейросети.

Проделано преимущественно на основе библиотеки Pytorch.

шаг 0: загрузка датасета(тренировка и тест), вывод примеров изображений

шаг 1: препроцессинг данных: разбиение на батчи по 1024 элемента(генератор), нормировка по первому батчу.

шаг 2: построение первой сети,

архитектура:

Conv: stride 3, pad 0, input channels 1, output channels 16

relu

Conv: stride 3, pad 0, input channels 16, output channels 64

relu

max pooling: kernel=2

flatten

relu

Linear

relu

Linear

Softmax

количество параметров 9216

шаг 3: тренировка: 5 эпох, оптимизатор Adam, learning rate = 0.001, потери - относительная кросс-энтропия

результат: потери 0.023

распределение ассурасу по цифрам:

0 accuracy: 0.9959183673469387

1 accuracy: 0.9964757709251101

2 accuracy: 0.9893410852713178

3 accuracy: 0.9871287128712871

4 accuracy: 0.9989816700610998

5 accuracy: 0.9854260089686099

6 accuracy: 0.9874739039665971

7 accuracy: 0.9785992217898832

8 accuracy: 0.9774127310061602

9 accuracy: 0.958374628344896

шаг 4: подсчет accuracy в среднем и для каждого класса, вывод примера предсказаний

результат: средняя accuracy 0.9855

шаг 5: построение второй сети на основе первой, добавлены batchnorm и dropout слои

шаг 6: тренировка (аналогично 3),

результат: потери 0.019

распределение accuracy по цифрам:

0 accuracy: 0.9979591836734694

1 accuracy: 0.9955947136563876

2 accuracy: 0.9903100775193798

3 accuracy: 0.9930693069306931

4 accuracy: 0.9928716904276986

5 accuracy: 0.992152466367713

6 accuracy: 0.9874739039665971

7 accuracy: 0.9863813229571985

8 accuracy: 0.9876796714579056

9 accuracy: 0.9821605550049554

шаг 7: аналогично 4

результат: средняя accuracy 0.9907

шаг 8: загрузка собственных цифр(10шт), приведение к датасетному виду, нормировка, оценка модели, вывод изображений с оценкой

результат: распознано 9/10

**Итог:** построен классификатор для распознавания рукописных цифр с точностью 0.9855

```

<import torchvision
import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision.transforms import ToTensor, Normalize, Compose

import warnings
warnings.filterwarnings("ignore")

dataset_train = torchvision.datasets.MNIST(root='/content/drive/MyDrive', download=False, transform=ToTensor)
dataset_test = torchvision.datasets.MNIST(root='/content/drive/MyDrive', train=False, download=False, transform=ToTensor)

X_train, X_test, y_train, y_test = dataset_train.data, dataset_test.data, dataset_train.targets, dataset_test.targets

num_test = X_test.shape[0]

import random
import numpy as np
import matplotlib.pyplot as plt
from IPython import display

%matplotlib inline
plt.rcParams['figure.figsize'] = (14.0, 12.0)
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

classes = list(np.arange(10))
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx])
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()>

```




Figure 1: This is a caption

```
<train_loader = torch.utils.data.DataLoader(dataset_train, shuffle=True, batch_size=1024, num_workers=1)
test_loader = torch.utils.data.DataLoader(dataset_test, shuffle=True, batch_size=1024)
```

```
data = next(iter(train_loader))
mean = data[0].mean()
std = data[0].std()
```

```
dataset_train = torchvision.datasets.MNIST(root='/content/drive/MyDrive', download=False, transform=Com
dataset_test = torchvision.datasets.MNIST(root='/content/drive/MyDrive', train=False, download=False, t
```

```
train_loader = torch.utils.data.DataLoader(dataset_train, shuffle=True, batch_size=1024, num_workers=1)
test_loader = torch.utils.data.DataLoader(dataset_test, shuffle=True, batch_size=1024)
```

```
device=torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
```

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 16, 3)
        self.conv2 = nn.Conv2d(16, 64, 3)
        self.lin1 = nn.Linear(64 * 12 * 12, 128)
        self.lin2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.conv1(x)
```

```

        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = torch.flatten(x, 1)
        x = self.lin1(x)
        x = F.relu(x)
        x = self.lin2(x)
        output = F.log_softmax(x, dim=1)
        return output

net = Net().to(device)

n_epoch = 5
loss_history = []

loss_fun = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(net.parameters(), lr=0.001)

for i in range(n_epoch):
    net.train()
    for batch, labels in train_loader:
        batch = batch.cuda()
        labels = labels.cuda()
        optimizer.zero_grad()

        pred = net.forward(batch)

        loss = loss_fun(pred, labels)
        loss.backward()

        optimizer.step()

        loss_history.append(loss)

    display.clear_output(wait=True)
    plt.figure(figsize=(8, 6))

    plt.title("Training loss")
    plt.xlabel("#iteration")
    plt.ylabel("loss")
    plt.plot(loss_history, 'b')
    plt.show()

    print('Current loss: %f' % loss)
>

```




Figure 2: This is a caption

```
<accs = torch.zeros(10)
losses = torch.zeros(10)
digit_shot = {i: 0 for i in range(10)}
digits_cnt = np.zeros(10)

net.eval()
with torch.no_grad():
    i = 0
    for batch, labels in test_loader:
        batch = batch.cuda()
        labels = labels.cuda()
        preds = net(batch)
        loss = loss_fun(preds, labels)
        label_preds = preds.argmax(dim=1)
        accuracy = (label_preds == labels).sum() / labels.shape[0]
        accs[i] = accuracy
        losses[i] = loss
        label_preds = label_preds.cpu().detach().numpy()
        labels = labels.cpu().detach().numpy()
        for j in range(len(labels)):
            if labels[j] == label_preds[j]:
                digit_shot[labels[j]] += 1
                digits_cnt[labels[j]] += 1
        i += 1

print("mean accuracy: ", accs.mean(),
      "mean loss: ", losses.mean())
```

```

for i in range(10):
    print("{} accuracy: {}".format(i, digit_shot[i] / digits_cnt[i]))>

mean accuracy: tensor(0.9855) mean loss: tensor(0.0472)
0 accuracy: 0.9959183673469387
1 accuracy: 0.9964757709251101
2 accuracy: 0.9893410852713178
3 accuracy: 0.9871287128712871
4 accuracy: 0.9989816700610998
5 accuracy: 0.9854260089686099
6 accuracy: 0.9874739039665971
7 accuracy: 0.9785992217898832
8 accuracy: 0.9774127310061602
9 accuracy: 0.958374628344896

mini_loader = torch.utils.data.DataLoader(dataset_test, shuffle=True, batch_size=100)
digit_preds = {i: [] for i in range(10)}
digits_cnt = np.zeros(10)

net.eval()
with torch.no_grad():
    batch, labels = next(iter(mini_loader))
    batch = batch.cuda()
    preds = net(batch)
    label_preds = preds.argmax(dim=1)
    label_preds = label_preds.cpu().detach().numpy()
    labels = labels.numpy()
    batch = batch.cpu().detach().numpy()
    for j in range(len(labels)):
        digit_preds[labels[j]].append(label_preds[j])
        digits_cnt[labels[j]] += 1

classes = list(np.arange(10))
num_classes = len(classes)
samples_per_class = int(min(digits_cnt))
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(label_preds.astype(np.int8) == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(batch[idx].reshape((28, 28)))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()

```




Figure 3: This is a caption

```
<class Net(nn.Module):
  def __init__(self):
    super(Net, self).__init__()
    self.conv1 = nn.Conv2d(1, 16, 3)
    self.bnorm1 = nn.BatchNorm2d(16)

    self.conv2 = nn.Conv2d(16, 64, 3)
    self.bnorm2 = nn.BatchNorm2d(64)
    self.dropout1 = nn.Dropout(0.3)

    self.lin1 = nn.Linear(64 * 12 * 12, 128)
    self.bnorm3 = nn.BatchNorm1d(128)
    self.lin2 = nn.Linear(128, 10)

  def forward(self, x):
    x = self.conv1(x)
    x = F.relu(x)
    x = self.bnorm1(x)
    x = F.relu(x)

    x = self.conv2(x)
    x = F.relu(x)
    x = self.bnorm2(x)
    x = F.relu(x)
    x = F.max_pool2d(x, 2)
    x = self.dropout1(x)
```



```

        x = torch.flatten(x, 1)
        x = self.lin1(x)
        x = self.bnorm3(x)
        x = F.relu(x)
        x = self.lin2(x)
        output = F.log_softmax(x, dim=1)
        return output

net = Net().to(device)

n_epoch = 5

loss_fun = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(net.parameters(), lr=0.001)

loss_history = []

for i in range(n_epoch):
    net.train()
    for batch, labels in train_loader:
        batch = batch.cuda()
        labels = labels.cuda()
        optimizer.zero_grad()

        pred = net.forward(batch)

        loss = loss_fun(pred, labels)
        loss_history.append(loss)
        loss.backward()

        optimizer.step()

    display.clear_output(wait=True)
    plt.figure(figsize=(8, 6))

    plt.title("Training loss")
    plt.xlabel("#iteration")
    plt.ylabel("loss")
    plt.plot(loss_history, 'b')
    plt.show()

    print('Current loss: %f' % loss)
>

```




Figure 4: This is a caption

```
<accs = torch.zeros(10)
losses = torch.zeros(10)
digit_shot = {i: 0 for i in range(10)}
digits_cnt = np.zeros(10)

net.eval()
with torch.no_grad():
    i = 0
    for batch, labels in test_loader:
        batch = batch.cuda()
        labels = labels.cuda()
        preds = net(batch)
        loss = loss_fun(preds, labels)
        label_preds = preds.argmax(dim=1)
        accuracy = (label_preds == labels).sum() / labels.shape[0]
        accs[i] = accuracy
        losses[i] = loss
        label_preds = label_preds.cpu().detach().numpy()
        labels = labels.cpu().detach().numpy()
        for j in range(len(labels)):
            if labels[j] == label_preds[j]:
                digit_shot[labels[j]] += 1
                digits_cnt[labels[j]] += 1
    i += 1

print("mean accuracy: ", accs.mean(),
      "mean loss: ", losses.mean())
```

```

for i in range(10):
    print("{} accuracy: {}".format(i, digit_shot[i] / digits_cnt[i]))

mean accuracy: tensor(0.9907) mean loss: tensor(0.0313)
0 accuracy: 0.9979591836734694
1 accuracy: 0.9955947136563876
2 accuracy: 0.9903100775193798
3 accuracy: 0.9930693069306931
4 accuracy: 0.9928716904276986
5 accuracy: 0.992152466367713
6 accuracy: 0.9874739039665971
7 accuracy: 0.9863813229571985
8 accuracy: 0.9876796714579056
9 accuracy: 0.9821605550049554

mini_loader = torch.utils.data.DataLoader(dataset_test, shuffle=True, batch_size=100)
digit_preds = {i: [] for i in range(10)}
digits_cnt = np.zeros(10)

net.eval()
with torch.no_grad():
    batch, labels = next(iter(mini_loader))
    batch = batch.cuda()
    preds = net(batch)
    label_preds = preds.argmax(dim=1)
    label_preds = label_preds.cpu().detach().numpy()
    labels = labels.numpy()
    batch = batch.cpu().detach().numpy()
    for j in range(len(labels)):
        digit_preds[labels[j]].append(label_preds[j])
        digits_cnt[labels[j]] += 1

classes = list(np.arange(10))
num_classes = len(classes)
samples_per_class = int(min(digits_cnt))
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(label_preds.astype(np.int8) == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(batch[idx].reshape((28, 28)))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()>

```




Figure 5: This is a caption

```
<mini_loader = torch.utils.data.DataLoader(dataset_test, shuffle=True, batch_size=1)
batch, labels = next(iter(mini_loader))

DIGITS_PATH = '/content/drive/MyDrive/digits_bold'

transform = torchvision.transforms.Compose([torchvision.transforms.Grayscale(1), torchvision.transforms
                                           torchvision.transforms.Resize((28,28)),
                                           torchvision.transforms.Lambda(lambda x: torchvision.transfo

my_digits = torchvision.datasets.ImageFolder(DIGITS_PATH, transform=transform)
mini_loader = torch.utils.data.DataLoader(my_digits, shuffle=True, batch_size=10)

data = next(iter(mini_loader))
mean = data[0].mean()
std = data[0].std()

transform = torchvision.transforms.Compose([torchvision.transforms.Grayscale(1), torchvision.transforms
                                           torchvision.transforms.Resize((28,28)),
                                           torchvision.transforms.Lambda(lambda x: torchvision.transfo
                                           Normalize(mean, std)])

my_digits = torchvision.datasets.ImageFolder(DIGITS_PATH, transform=transform)

mini_loader = torch.utils.data.DataLoader(my_digits, shuffle=True, batch_size=10)
digit_preds = {i: [] for i in range(10)}
digits_cnt = np.zeros(10)

net.eval()
```

```

with torch.no_grad():
    batch, labels = next(iter(mini_loader))
    batch = batch.cuda()
    preds = net(batch)
    label_preds = preds.argmax(dim=1)
    label_preds = label_preds.cpu().detach().numpy()
    labels = labels.numpy()
    batch = batch.cpu().detach().numpy()
    for j in range(len(labels)):
        digit_preds[labels[j]].append(label_preds[j])
        digits_cnt[labels[j]] += 1

label_preds

array([7, 7, 0, 2, 3, 1, 6, 8, 4, 5])

classes = list(np.arange(10))
num_classes = len(classes)
samples_per_class = 1
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(labels.astype(np.int8) == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(batch[idx].reshape((28, 28)))
        plt.axis('off')
        if i == 0:
            plt.title(label_preds[idx])
plt.show()>

```



Figure 6: This is a caption